Generative AI on Google Cloud

Table of Contents

1. Introduction
    - Overview of Generative AI
    - Why deploy Generative AI on the cloud?
    - Benefits of cloud-based AI models
2. Choosing a Cloud Provider
    - AWS (Amazon Bedrock, SageMaker)
    - Google Cloud (Vertex AI, AI Platform)
    - Microsoft Azure (Azure OpenAI, ML Studio)
    - Hugging Face Spaces
    - IBM Watson AI
3. Setting Up the Cloud Environment
    - Creating an account and setting up billing
    - Configuring GPU/TPU instances
    - Installing necessary libraries (TensorFlow, PyTorch, Transformers, etc.)
4. Training and Fine-Tuning Models
    - Choosing a pre-trained model (Stable Diffusion, LLaMA, GPT, etc.)
    - Uploading datasets and preprocessing
    - Fine-tuning the model using cloud-based GPUs/TPUs
5. Deploying the Model
    - Using APIs for inference (FastAPI, Flask, Streamlit)
    - Serverless deployment with AWS Lambda, Google Cloud Functions
    - Scaling with Kubernetes, Docker
6. Monitoring and Optimization
    - Logging model performance
    - Reducing cloud costs
    - Security considerations
7. Use Cases and Applications
    - Text generation, Image generation, Code generation
    - Real-world applications

8. Conclusion and Next Steps
    - Future trends in cloud-based Generative AI
    - Learning resources
9. Interview Question

1. Introduction

What is Generative AI?

Generative AI refers to artificial intelligence models that can generate text, images, audio, or other data based on training inputs. These models include GPT (for text), Stable Diffusion (for images), and music generation models like Jukebox.

Why Use Cloud for Generative AI?

- **Scalability**: Cloud platforms offer on-demand computing power, including GPUs and TPUs.
- **Cost-effectiveness**: Pay-as-you-go pricing models reduce hardware costs.
- **Easy Deployment**: Cloud providers offer managed services for quick model deployment.
- **Collaboration**: Teams can work together remotely on AI projects.

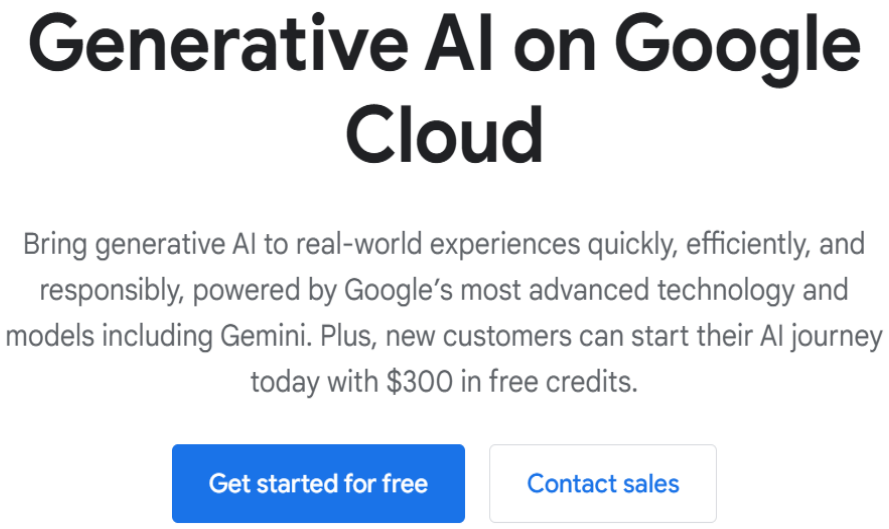

The image is promoting **Generative AI on Google Cloud**, highlighting its efficiency, responsibility, and advanced models like **Gemini**. It mentions that new customers can start with **$300 in free credits** and provides two call-to-action buttons:

- **"Get started for free"** (likely leading to a signup page).
- **"Contact sales"** (probably for enterprise inquiries).

2. Choosing a Cloud Provider

Different cloud platforms provide various AI tools and services. Below are some of the best options:

AWS (Amazon Web Services)

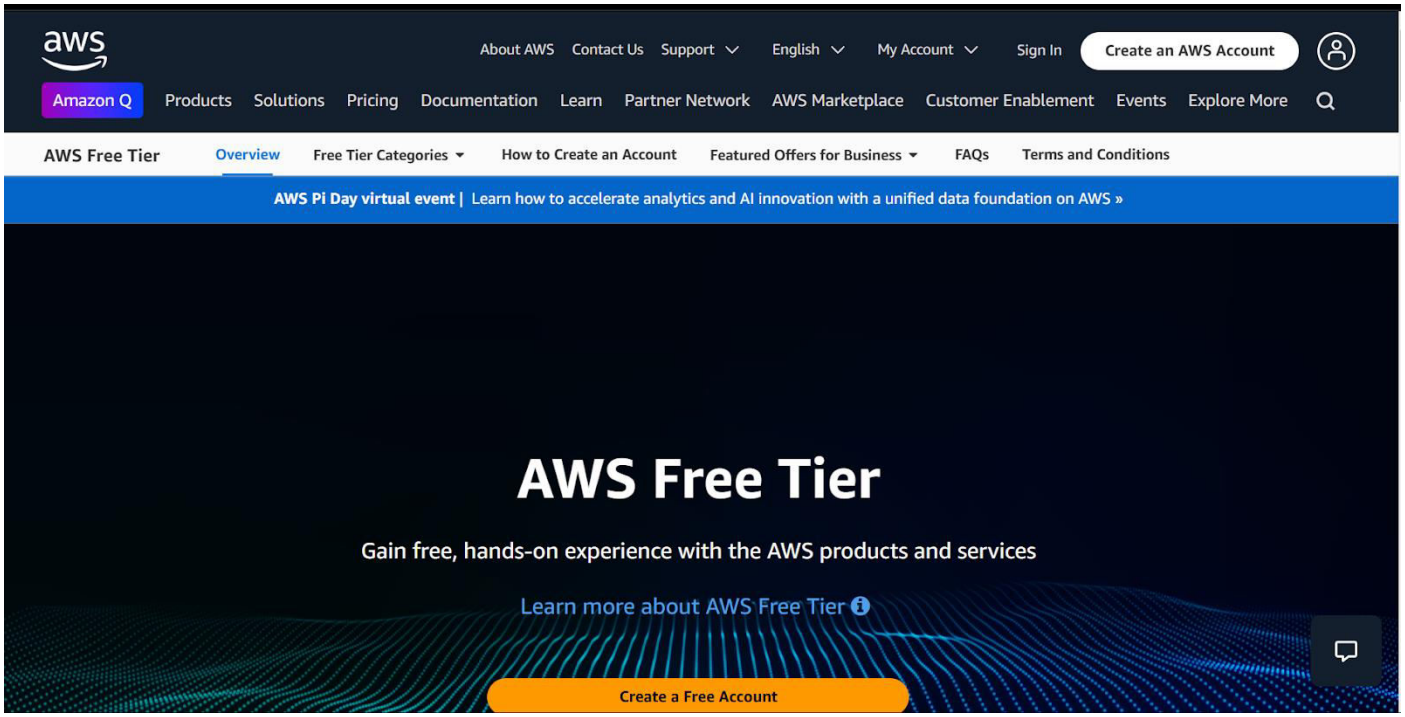

- **Amazon SageMaker**: Train and deploy AI models.
- **AWS Lambda**: Serverless execution for inference.
- **Amazon Bedrock**: Pre-trained foundation models for easy integration.

Google Cloud

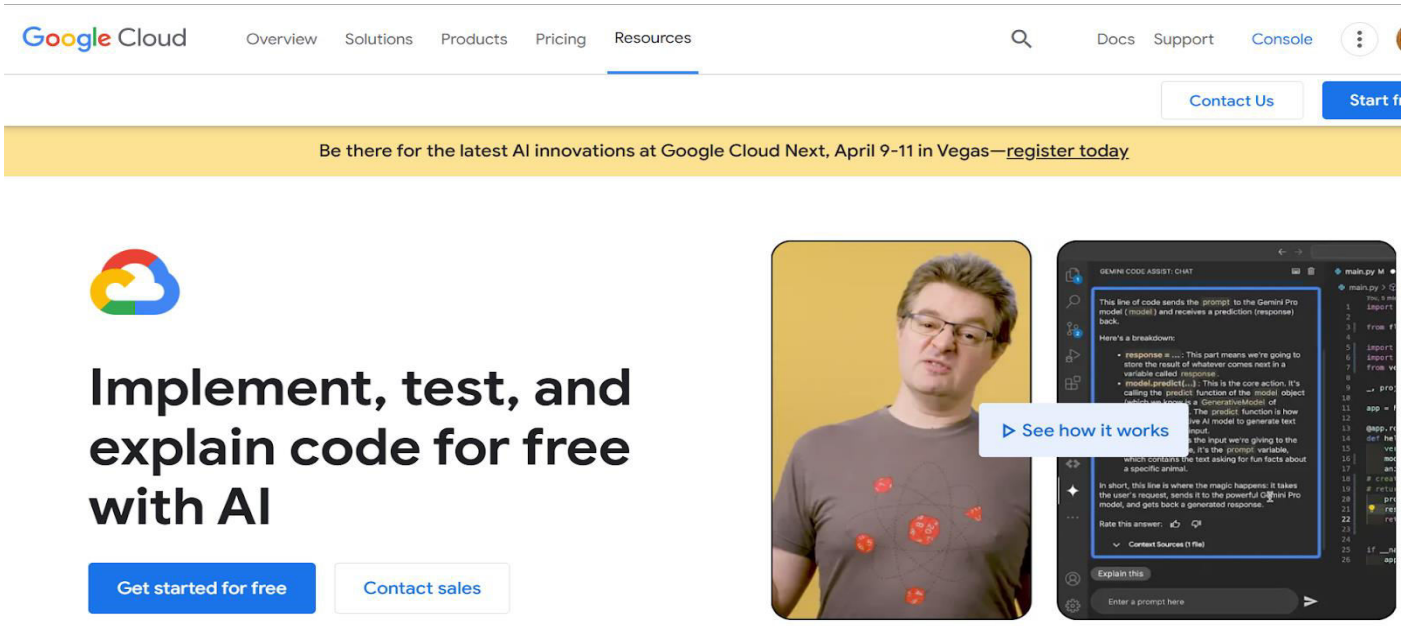

- **Vertex AI**: A unified AI platform for training and deployment.
- **Google Cloud TPU**: Specialized hardware for deep learning.
- **BigQuery ML**: Train ML models directly on structured data.

Microsoft Azure

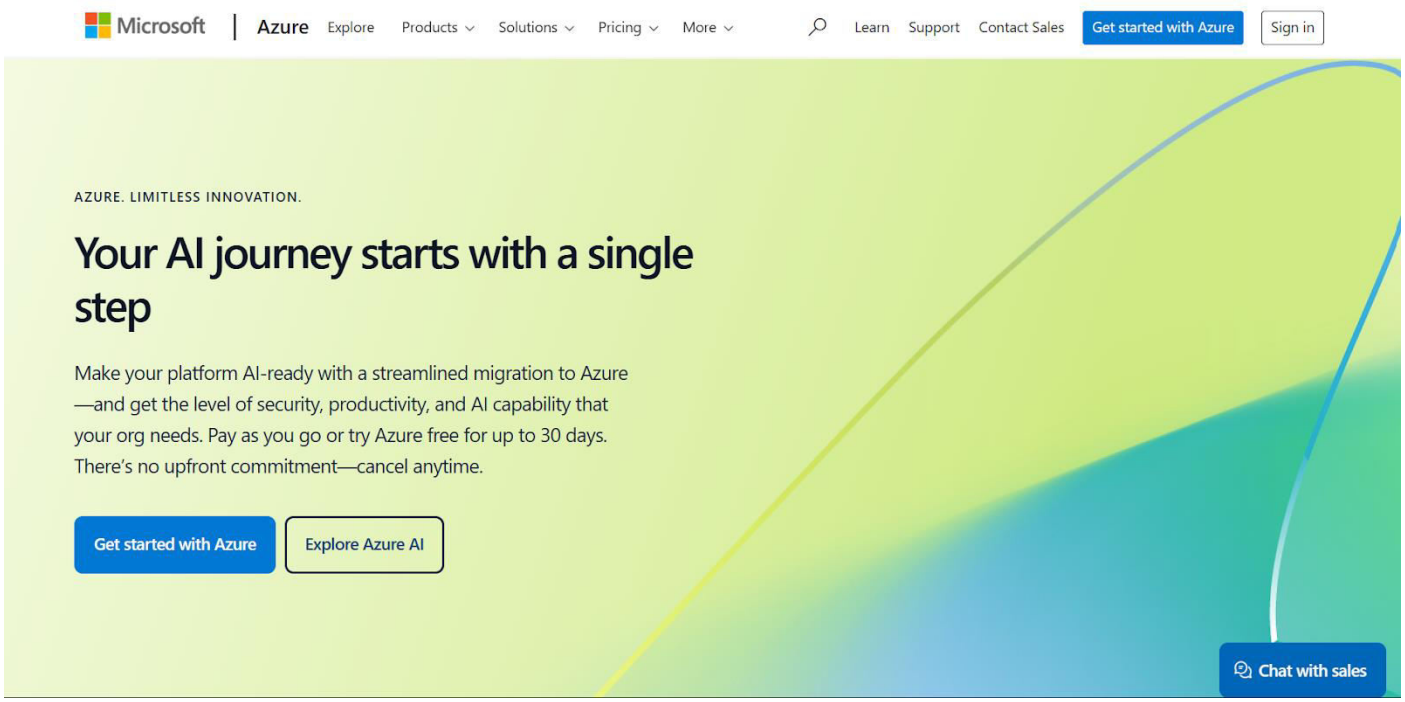

- **Azure OpenAI Service**: Access OpenAI's models like GPT-4.
- **Azure ML Studio**: Drag-and-drop interface for building ML models.
- **Azure Functions**: Serverless compute for AI workflows.

Other Platforms

- **Hugging Face Spaces**: Host AI models and demos easily.

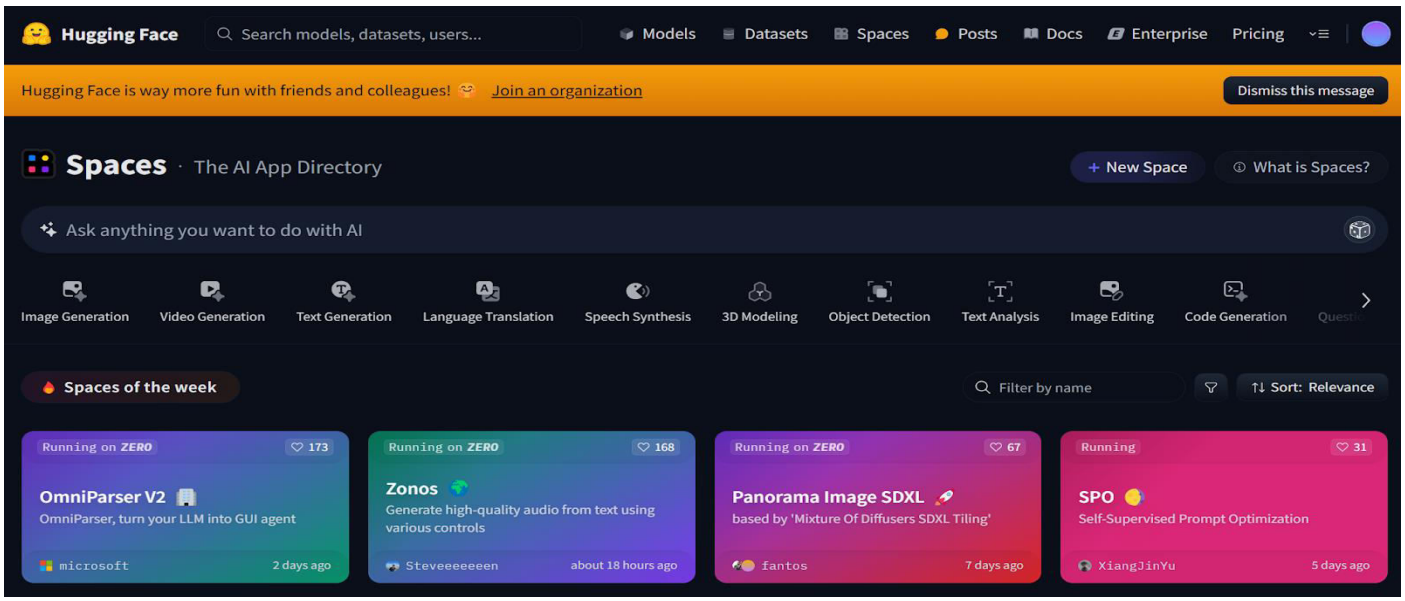

- **IBM Watson AI**: Enterprise AI solutions for NLP and data analysis.

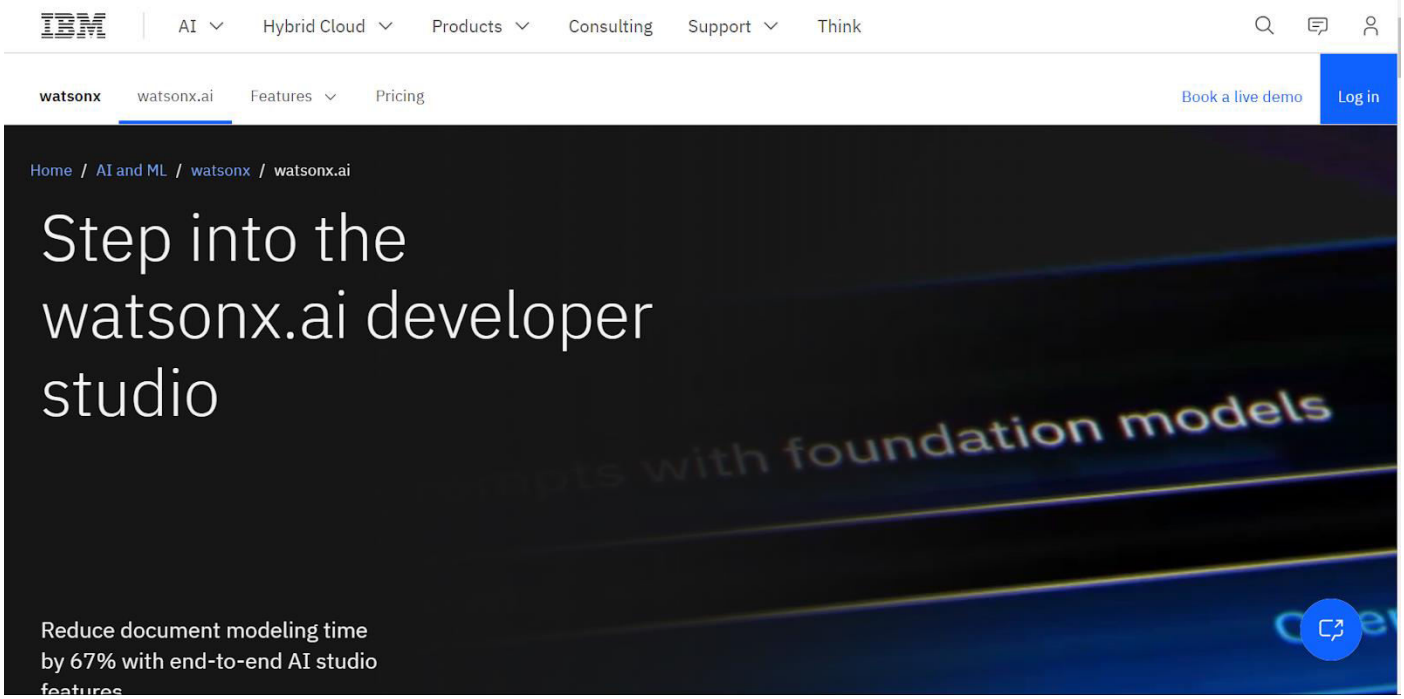

3. Setting Up the Cloud Environment

Step 1: Create an Account

1. Sign up for an account on AWS, Google Cloud, or Azure.
2. Enable billing and set up free-tier usage.

Step 2: Set Up Compute Resources

- Choose **GPU instances** (e.g., NVIDIA A100,V100) or **TPU instances** for better performance.
- Use **AutoML tools** for simplified training.

Step 3: Install Required Libraries

```
pip install torch tensorflow transformers datasets
```

- Install cloud SDKs:

```
# AWS
pip install boto3
# Google Cloud
pip install google-cloud-storage
# Azure
pip install azureml-sdk
```

4. Training and Fine-Tuning Models

Using Pre-trained Models

Hugging Face provides pre-trained models that can be fine-tuned on custom datasets:

```
from transformers import pipeline
model = pipeline("text-generation", model="gpt-3.5-turbo")
print(model("Hello, how are you?"))
```

Fine-Tuning a Model on Custom Data

1. Upload dataset to cloud storage.
2. Use cloud training services (SageMaker, Vertex AI, Azure ML).
3. Train the model using TPUs/GPUs.

5. Deploying the Model

Option 1: Deploy as API using FastAPI

```
from fastapi import FastAPI
from transformers import pipeline

app = FastAPI()
model = pipeline("text-generation", model="gpt-3.5-turbo")

@app.get("/generate")
def generate_text(prompt: str):
    return model(prompt)
```

- Deploy on AWS Lambda, Google Cloud Functions, or Azure Functions.

Option 2: Deploy as a Web App using Streamlit

```
import streamlit as st
st.title("Generative AI Demo")
prompt = st.text_input("Enter text:")
if prompt:
    st.write(model(prompt))
```

6. Monitoring and Optimization

Logging Model Performance

Use MLflow or TensorBoard to track performance.

```
pip install mlflow
docker run -p 5000:5000 mlflow/mlflow
```

Optimizing for Cost & Efficiency

- Use **Spot Instances** (AWS) or **Preemptible VMs** (GCP) for lower costs.
- Deploy using **serverless architectures** to reduce idle compute costs.
- Optimize model size using **quantization** and **pruning**.

7. Use Cases and Applications

- **Text Generation** (Chatbots, Content Creation)
- **Image Generation** (Stable Diffusion, DALL·E)
- **Code Generation** (Codex, Copilot)
- **Audio Generation** (AI Music, Speech Synthesis)

8. Conclusion & Next Steps

- Learn more about **AutoML** and **Hyperparameter Tuning**.
- Explore **Multi-Cloud Strategies** for better reliability.
- Follow **latest advancements** in AI and cloud technologies.

Resources

- AWS Machine Learning
- Google Vertex AI
- Azure Machine Learning
- Hugging Face Model Hub

Interview Questions

Basic Level Questions

1. What is Generative AI, and how does it differ from traditional AI?
Generative AI creates new content (text, images, audio, etc.) using deep learning models, while **traditional AI** focuses on pattern recognition and predictions.

2. Why should we use cloud platforms for building Generative AI models?
Cloud platforms provide **high-performance GPUs/TPUs**, **scalability**, **cost efficiency**, and **easy deployment** compared to local hardware.

3. What are the advantages of using cloud GPUs/TPUs over local hardware?
No upfront hardware cost
Access to powerful AI-optimized processors
Faster training and inference
Scalability (adjust resources as needed)

4. Explain the role of AutoML in training AI models on the cloud.
AutoML automates the process of **model selection, training, hyperparameter tuning, and deployment**, making AI accessible to non-experts.

5. What are the key differences between AWS SageMaker, Google Vertex AI, and Azure Machine Learning?

Feature	AWS SageMaker	Google Vertex AI	Azure ML
Compute	EC2, GPU, TPU	TPU, AI Platform	GPU, ML Studio
Deployment	API, Lambda	Cloud Run	Azure Functions
Model Training	AutoML, Custom	AutoML, TensorFlow	AutoML, Designer

6. What is the purpose of Hugging Face in Generative AI?
Hugging Face provides **pre-trained models (GPT, BERT, Stable Diffusion)** and tools for **fine-tuning and deploying models** on cloud platforms.

7. How do you choose between serverless and containerized deployment for an AI model?
Use Serverless (AWS Lambda, Cloud Functions): For lightweight, event-driven AI tasks.
Use Containerized Deployment (Docker, Kubernetes): For large-scale, persistent AI services.

8. What are the common challenges faced while deploying AI models on the cloud?
Model latency
Cost optimization
Security & data privacy
Managing large datasets

9. What is MLOps, and how does it help in managing cloud AI deployments?
MLOps (Machine Learning Operations) automates the AI lifecycle: **model versioning, CI/CD pipelines, monitoring, and retraining** for cloud-based AI.

10. Explain the importance of data preprocessing and augmentation before training a model.
Preprocessing (cleaning, normalizing) improves model accuracy.
Data augmentation (rotation, flipping, noise addition) prevents overfitting and enhances generalization.

Intermediate Level Questions

11. How do you fine-tune a pre-trained Transformer model on cloud infrastructure?
Choose a **pre-trained model** (e.g., GPT-3, BERT from Hugging Face).
Upload **custom dataset** to cloud storage (S3, GCS, Azure Blob).
Use cloud **compute instances (GPU/TPU)** to fine-tune the model.
Deploy the model via **API, container, or serverless service**.

12. What are the differences between training and inference workloads in cloud-based AI?
Training: Requires high-power GPUs/TPUs for large datasets.
Inference: Uses optimized, smaller compute resources for real-time predictions.

13. How does model quantization help optimize AI models for cloud deployment?
Model quantization reduces precision (e.g., from **FP32 to INT8**), making models:
Smaller in size

Faster in inference
More efficient in low-resource environments

14. What are Spot Instances (AWS) and Preemptible VMs (GCP), and how do they help reduce costs?
Spot Instances (AWS): Up to **90% cheaper**, but can be interrupted anytime.
Preemptible VMs (GCP): Low-cost, short-term instances useful for batch AI jobs.

15. How do you monitor an AI model's performance on the cloud after deployment?
Use **MLflow** or **TensorBoard** for logging model metrics.
Implement **CloudWatch (AWS)** or **Stackdriver (Google Cloud)** for real-time monitoring.

16. What is Feature Store, and how does it help in Generative AI model training?
A **Feature Store** is a **centralized repository** for storing and managing features across multiple AI models, ensuring **consistency and reusability**.

17. How do you secure sensitive data when training AI models on a public cloud?
Encrypt data (at rest and in transit).
Use **IAM (Identity & Access Management)**.
Implement **VPC & Private Networking** for secure access.

18. What are the best practices for scaling AI workloads on cloud platforms?
Use **auto-scaling** for dynamic resource allocation.
Optimize **batch processing** for large datasets.
Deploy AI models via **Kubernetes for distributed workloads**.

19. How can APIs like OpenAI API, Gemini API, or Claude API be integrated into cloud-based applications?
Use **REST API endpoints** to call AI models.
Authenticate using **API keys** or **OAuth tokens**.
Deploy AI-based chatbots, automation tools, and image generators.

20. What are some real-world applications of Generative AI on Cloud in different industries?
Healthcare: AI-driven diagnosis, medical image synthesis.
Finance: Fraud detection, AI-generated investment reports.
E-commerce: Product recommendations, AI-powered search.
Entertainment: AI-generated music, movies, and deepfake detection.

THANKYOU